

Project Name : - Algerian Forest Fire Dataset EDA , FE & Temperature Prediction using Linear Regression, Ridge, Lasso, Elastic Net .

1) Problem statement .

- This dataset comprises of Algerian Forest Fire Dataset taken from UCI .
- Link of the dataset is as follows :- <https://archive.ics.uci.edu/ml/datasets/Algerian+Forest+Fires+Dataset++> .

2) Data Collection.

- This dataset includes 244 instances that regroup a data of 2 regions of Algeria, namely the Brjajia region located in the northeast of Algeria and the Sidi Bel-abbes region located in the northwest of Algeria .
- 122 instances for each region .
- The Period is from June 2012 to September 2012. The Dataset includes 11 attributes and 1 output attribute i.e. Temperature
- The data consists of 14 column and 246 rows.

2.1 Import Data and Required Packages

Importing Necessary Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.linear_model import ElasticNet
from sklearn.linear_model import Lasso
from sklearn.linear_model import Ridge
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_absolute_error
from warnings import filterwarnings
filterwarnings('ignore')
%matplotlib inline
```

Loading the Algerian Forest Fire Dataset

```
In [2]: df=pd.read_csv("Algerian_forest_fires_dataset_UPDATE.csv",header=1)
```

```
In [3]: df.head()
```

```
Out[3]:
```

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
0	01	06	2012	29	57	18	0	65.7	3.4	7.6	1.3	3.4	0.5	not fire
1	02	06	2012	29	61	13	1.3	64.4	4.1	7.6	1	3.9	0.4	not fire
2	03	06	2012	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	not fire

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
3	04	06	2012	25	89	13	2.5	28.6	1.3	6.9	0	1.7	0	not fire
4	05	06	2012	27	77	16	0	64.8	3	14.2	1.2	3.9	0.5	not fire

Attribute Information :-

Period Covered

- 1. Date : (DD/MM/YYYY) Day, month ('june' to 'september'), year (2012)

Weather data observations

- 1. Temp : temperature noon (temperature max) in Celsius degrees: 22 to 42
- 1. RH : Relative Humidity in %: 21 to 90
- 1. Ws :Wind speed in km/h: 6 to 29
- 1. Rain: total day in mm: 0 to 16.8

FWI Components

- 1. Fine Fuel Moisture Code (FFMC) index from the FWI system: 28.6 to 92.5
- 1. Duff Moisture Code (DMC) index from the FWI system: 1.1 to 65.9
- 1. Drought Code (DC) index from the FWI system: 7 to 220.4
- 1. Initial Spread Index (ISI) index from the FWI system: 0 to 18.5
- 1. Buildup Index (BUI) index from the FWI system: 1.1 to 68
- 1. Fire Weather Index (FWI) Index: 0 to 31.1
- 1. Classes: two classes, namely "Fire" and "not Fire"

In [4]: `df.tail()`

Out[4]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
241	26	09	2012	30	65	14	0	85.4	16	44.5	4.5	16.9	6.5	fire
242	27	09	2012	28	87	15	4.4	41.1	6.5	8	0.1	6.2	0	not fire
243	28	09	2012	27	87	29	0.5	45.9	3.5	7.9	0.4	3.4	0.2	not fire
244	29	09	2012	24	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7	not fire
245	30	09	2012	24	64	15	0.2	67.3	3.8	16.5	1.2	4.8	0.5	not fire

In [5]: `df.shape`

Out[5]: (246, 14)

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 246 entries, 0 to 245
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   day         246 non-null   object
1   month       245 non-null   object
2   year        245 non-null   object
```

```
3   Temperature    245 non-null    object
4   RH             245 non-null    object
5   Ws             245 non-null    object
6   Rain           245 non-null    object
7   FFMC           245 non-null    object
8   DMC            245 non-null    object
9   DC             245 non-null    object
10  ISI            245 non-null    object
11  BUI            245 non-null    object
12  FWI            245 non-null    object
13  Classes        244 non-null    object
```

dtypes: object(14)
memory usage: 27.0+ KB

```
In [7]: df.isnull().sum()
```

```
Out[7]: day                0
month                1
year                1
Temperature          1
RH                  1
Ws                  1
Rain                1
FFMC                1
DMC                 1
DC                  1
ISI                 1
BUI                 1
FWI                 1
Classes             2
dtype: int64
```

```
In [8]: df.describe()
```

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
count	246	245	245	245	245	245	245	245	245	245	245	245	245	244
unique	33	5	2	20	63	19	40	174	167	199	107	175	128	9
top	01	07	2012	35	64	14	0	88.9	7.9	8	1.1	3	0.4	fire
freq	8	62	244	29	10	43	133	8	5	5	8	5	12	131

```
In [9]: df.iloc[121:125,:]
```

Out[9]:

	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
121	30	09	2012	25	78	14	1.4	45	1.9	7.5	0.2	2.4	0.1	not fire
122	Sidi-Bel Abbes Region Dataset	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN
123	day	month	year	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes
124	01	06	2012	32	71	12	0.7	57.1	2.5	8.2	0.6	2.8	0.2	not fire

```
In [ ]:
```

```
In [ ]:
```

2.2 Data Cleaning

Dropping row no 122 specifying region name & 123 respecifying the header

```
In [10]: df.drop([122,123],inplace=True)
```

Resetting the index and dropping the index column

```
In [11]: df.reset_index(inplace=True)
df.drop('index',axis=1,inplace=True)
```

```
In [ ]:
```

Creating a new column called Region representing [0:- Bejaia and 1- Sidi Bel-abbes]

```
In [12]: df.loc[:122,"Region"]=0
df.loc[122:,"Region"]=1
```

Checking the Column Headers

```
In [13]: df.columns
```

```
Out[13]: Index(['day', 'month', 'year', 'Temperature', 'RH', 'Ws', 'Rain ', 'FFMC',
               'DMC', 'DC', 'ISI', 'BUI', 'FWI', 'Classes ', 'Region'],
              dtype='object')
```

Removing unnecessary space in column headers using str.strip()

```
In [14]: df.columns=df.columns.str.strip()
df.columns
```

```
Out[14]: Index(['day', 'month', 'year', 'Temperature', 'RH', 'Ws', 'Rain', 'FFMC',
               'DMC', 'DC', 'ISI', 'BUI', 'FWI', 'Classes', 'Region'],
              dtype='object')
```

Dropping rows with null values

```
In [15]: df.dropna(inplace=True)
```

Converting the necessary column dataye to int

```
In [16]: df.dtypes
```

```
Out[16]: day          object
month         object
year          object
Temperature    object
RH            object
Ws            object
Rain          object
FFMC          object
DMC           object
DC            object
```

```
ISI                object
BUI                object
FWI                object
Classes            object
Region             float64
dtype: object
```

```
In [17]: df[['day', 'month', 'year', 'Temperature', 'RH', 'Ws', "Region"]]=df[['day', 'month', 'year', 'Temperature', 'RH', 'Ws', 'Region']]
```

```
In [18]: df.dtypes
```

```
Out[18]: day                int32
month              int32
year              int32
Temperature        int32
RH                 int32
Ws                 int32
Rain              object
FFMC               object
DMC               object
DC                object
ISI               object
BUI               object
FWI               object
Classes            object
Region            int32
dtype: object
```

Values in df[Classes] has unnecessary spaces that are removed by str.strip()

```
In [19]: df.Classes.unique()
```

```
Out[19]: array(['not fire   ', 'fire   ', 'fire', 'fire ', 'not fire', 'not fire ',
               'not fire   ', 'not fire   '], dtype=object)
```

```
In [20]: df.Classes=df.Classes.str.strip()
df.Classes.unique()
```

```
Out[20]: array(['not fire', 'fire'], dtype=object)
```

Converting the Necessary Column Datatype to Float

```
In [21]: df.columns
```

```
Out[21]: Index(['day', 'month', 'year', 'Temperature', 'RH', 'Ws', 'Rain', 'FFMC',
               'DMC', 'DC', 'ISI', 'BUI', 'FWI', 'Classes', 'Region'],
              dtype='object')
```

```
In [22]: df[['Rain', 'FFMC', 'DMC', 'DC', 'ISI', 'BUI', 'FWI']]=df[['Rain', 'FFMC', 'DMC', 'DC', 'ISI', 'BUI', 'FWI']].astype(float)
```

```
In [23]: df.dtypes
```

```
Out[23]: day                int32
month              int32
year              int32
Temperature        int32
RH                 int32
Ws                 int32
Rain              float64
```

```
FFMC      float64
DMC        float64
DC          float64
ISI         float64
BUI         float64
FWI         float64
Classes    object
Region     int32
dtype: object
```

Dropping the year column as the data is for the same year

```
In [24]: df1=df.drop(['year'],axis=1)
```

DataFrame Description

```
In [25]: df1.describe().T
```

```
Out[25]:
```

	count	mean	std	min	25%	50%	75%	max
day	243.0	15.761317	8.842552	1.0	8.00	16.0	23.00	31.0
month	243.0	7.502058	1.114793	6.0	7.00	8.0	8.00	9.0
Temperature	243.0	32.152263	3.628039	22.0	30.00	32.0	35.00	42.0
RH	243.0	62.041152	14.828160	21.0	52.50	63.0	73.50	90.0
Ws	243.0	15.493827	2.811385	6.0	14.00	15.0	17.00	29.0
Rain	243.0	0.762963	2.003207	0.0	0.00	0.0	0.50	16.8
FFMC	243.0	77.842387	14.349641	28.6	71.85	83.3	88.30	96.0
DMC	243.0	14.680658	12.393040	0.7	5.80	11.3	20.80	65.9
DC	243.0	49.430864	47.665606	6.9	12.35	33.1	69.10	220.4
ISI	243.0	4.742387	4.154234	0.0	1.40	3.5	7.25	19.0
BUI	243.0	16.690535	14.228421	1.1	6.00	12.4	22.65	68.0
FWI	243.0	7.035391	7.440568	0.0	0.70	4.2	11.45	31.1
Region	243.0	0.497942	0.501028	0.0	0.00	0.0	1.00	1.0

3. Exploratory Data Analysis

Encoding not fire as 0 and Fire as 1

```
In [26]: df1['Classes']=np.where(df1['Classes']=='not fire ',0,1)
```

```
In [27]: df1.head()
```

```
Out[27]:
```

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	Region
0	1	6	29	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	1	0
1	2	6	29	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	1	0
2	3	6	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	1	0

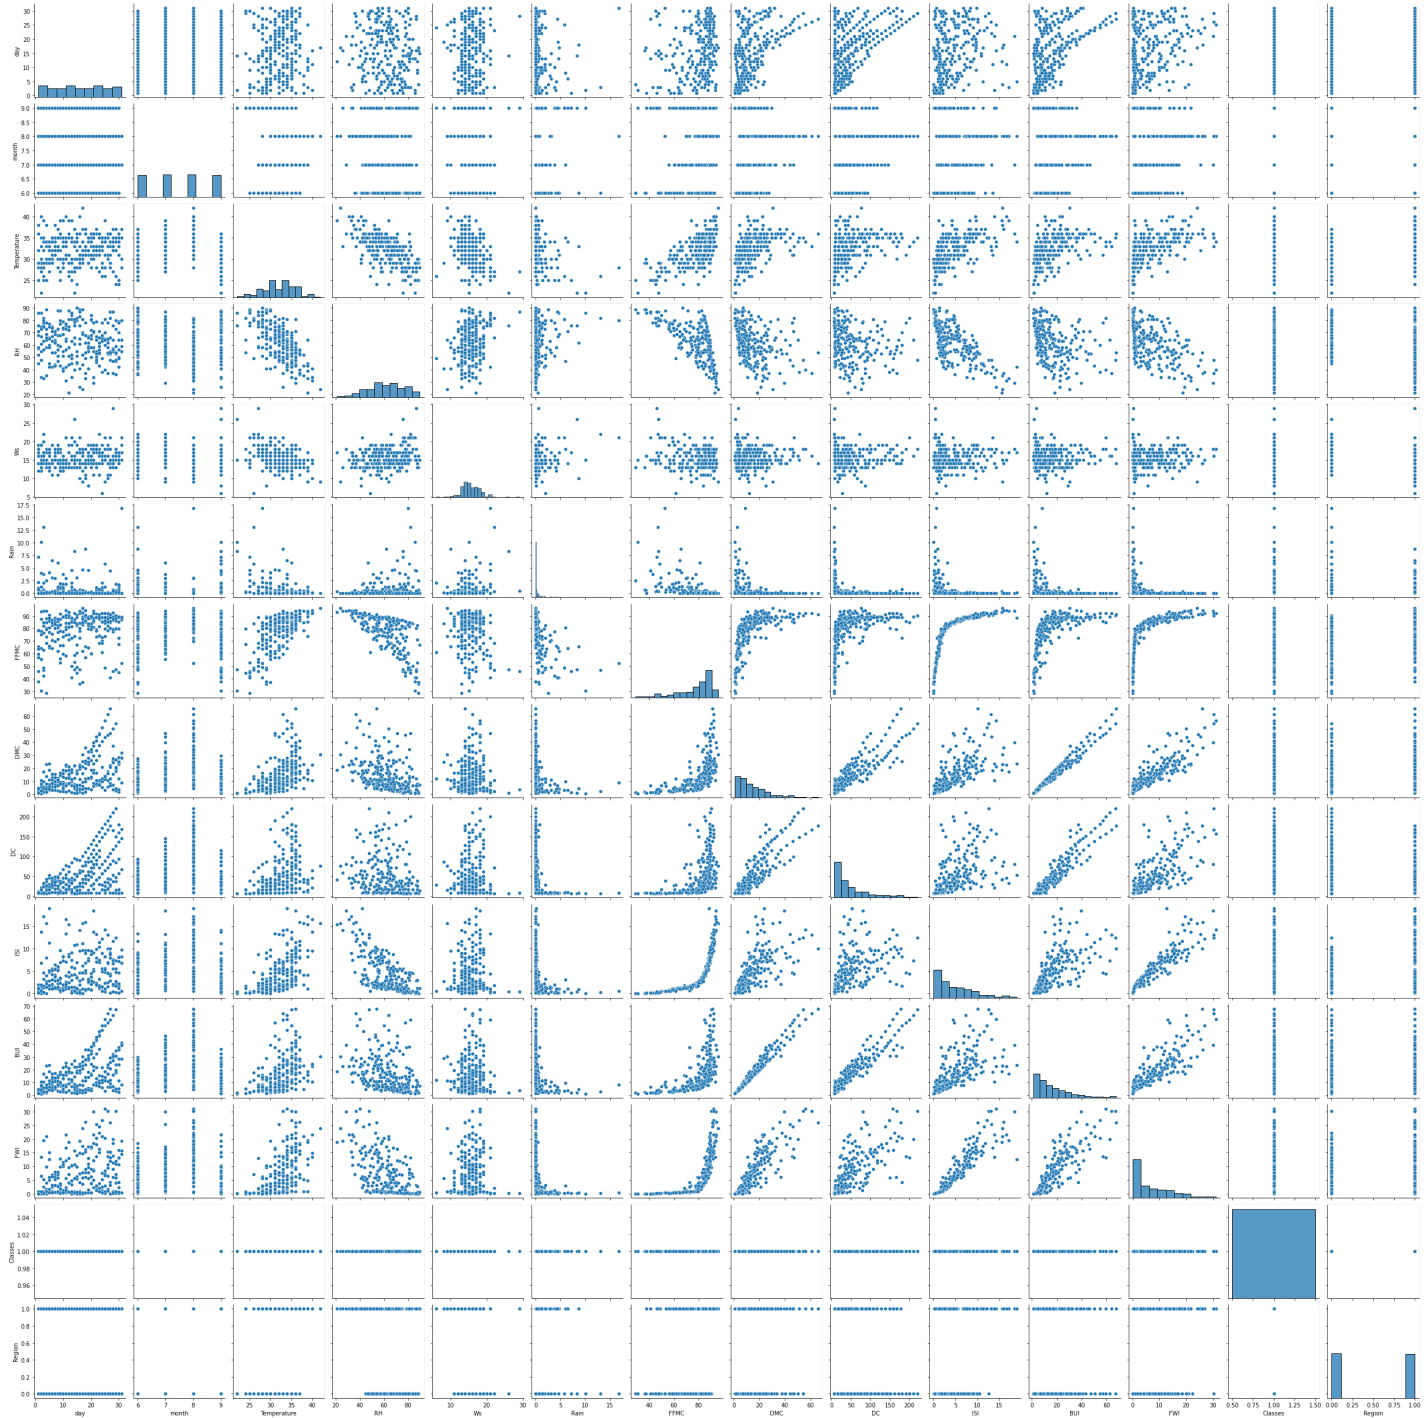
	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	Region
3	4	6	25	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0.0	1	0
4	5	6	27	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	1	0

```
In [28]: df1.corr()
```

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	Region
day	1.000000	-0.000369	0.097227	-0.076034	0.047812	-0.112523	0.224956	0.491514	0.527952	0.180543	0.065608	0.350781	0.000821	0.001857
month	-0.000369	1.000000	-0.056781	-0.041252	-0.039880	0.034822	0.017030	0.067943	0.126511	0.065608	1.000000	-0.580957	-0.402682	-0.181160
Temperature	0.097227	-0.056781	1.000000	-0.651400	-0.284510	-0.326492	0.676568	0.485687	0.376284	0.603871	-0.686667	0.032368	-0.580957	0.032368
RH	-0.076034	-0.041252	-0.651400	1.000000	0.244048	0.222356	-0.644873	-0.408519	-0.226941	0.079135	-0.298023	0.507397	0.875925	1.000000
Ws	0.047812	-0.039880	-0.284510	0.244048	1.000000	0.171506	-0.166548	-0.000721	0.079135	-0.298023	0.507397	0.875925	1.000000	0.000000
Rain	-0.112523	0.034822	-0.326492	0.222356	0.171506	1.000000	-0.543906	-0.288773	-0.298023	0.507397	0.875925	1.000000	0.000000	0.000000
FFMC	0.224956	0.017030	0.676568	-0.644873	-0.166548	-0.543906	1.000000	0.603608	0.507397	0.875925	1.000000	0.000000	0.000000	0.000000
DMC	0.491514	0.067943	0.485687	-0.408519	-0.000721	-0.288773	0.603608	1.000000	0.875925	1.000000	0.000000	0.000000	0.000000	0.000000
DC	0.527952	0.126511	0.376284	-0.226941	0.079135	-0.298023	0.507397	0.875925	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000
ISI	0.180543	0.065608	0.603871	-0.686667	0.008532	-0.347484	0.740007	0.680454	0.508643	1.000000	0.000000	0.000000	0.000000	0.000000
BUI	0.517117	0.085073	0.459789	-0.353841	0.031438	-0.299852	0.592011	0.982248	0.941988	0.000000	1.000000	0.000000	0.000000	0.000000
FWI	0.350781	0.082639	0.566670	-0.580957	0.032368	-0.324422	0.691132	0.875864	0.739521	0.000000	0.000000	1.000000	0.000000	0.000000
Classes	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	1.000000	0.000000
Region	0.000821	0.001857	0.269555	-0.402682	-0.181160	-0.040013	0.222241	0.192089	-0.078734	0.000000	0.000000	0.000000	0.000000	1.000000


```
In [29]: sns.pairplot(df1)
```

```
Out[29]: <seaborn.axisgrid.PairGrid at 0x1e875ab4190>
```



```
In [30]: sns.set(rc={'figure.figsize':(15,10)})
sns.heatmap(df1.corr(),annot=True)
```

```
Out[30]: <AxesSubplot:>
```

Report


- RH is negatively correlated with Temperature, FFMFC and ISI
- Rain is negatively correlated with Temperature and FFMFC, DMC, ISI and BUI

```
In [31]: df1.columns
```

```
Out[31]: Index(['day', 'month', 'Temperature', 'RH', 'Ws', 'Rain', 'FFMC', 'DMC', 'DC',
        'ISI', 'BUI', 'FWI', 'Classes', 'Region'],
        dtype='object')
```


```
In [32]: plt.scatter(df1['FFMC'], df1['Temperature'])
plt.xlabel("FFMC")
plt.ylabel("Temperature")
```

```
Out[32]: Text(0, 0.5, 'Temperature')
```




```
In [33]: sns.regplot(x='FFMC',y='Temperature',data=df1)
```

```
Out[33]: <AxesSubplot:xlabel='FFMC', ylabel='Temperature'>
```




```
In [34]: plt.scatter(df1['FWI'],df1['Temperature'])  
plt.xlabel("FWI")  
plt.ylabel("Temperature")
```

```
Out[34]: Text(0, 0.5, 'Temperature')
```




```
In [35]: sns.regplot(x='FWI',y='Temperature',data=df1)
```

```
Out[35]: <AxesSubplot:xlabel='FWI', ylabel='Temperature'>
```



```
In [36]: plt.scatter(df1['DMC'],df1['Temperature'])  
plt.xlabel("DMC")  
plt.ylabel("Temperature")
```

```
Out[36]: Text(0, 0.5, 'Temperature')
```



```
In [37]: sns.regplot(x='DMC',y='Temperature',data=df1)
```

```
Out[37]: <AxesSubplot:xlabel='DMC', ylabel='Temperature'>
```

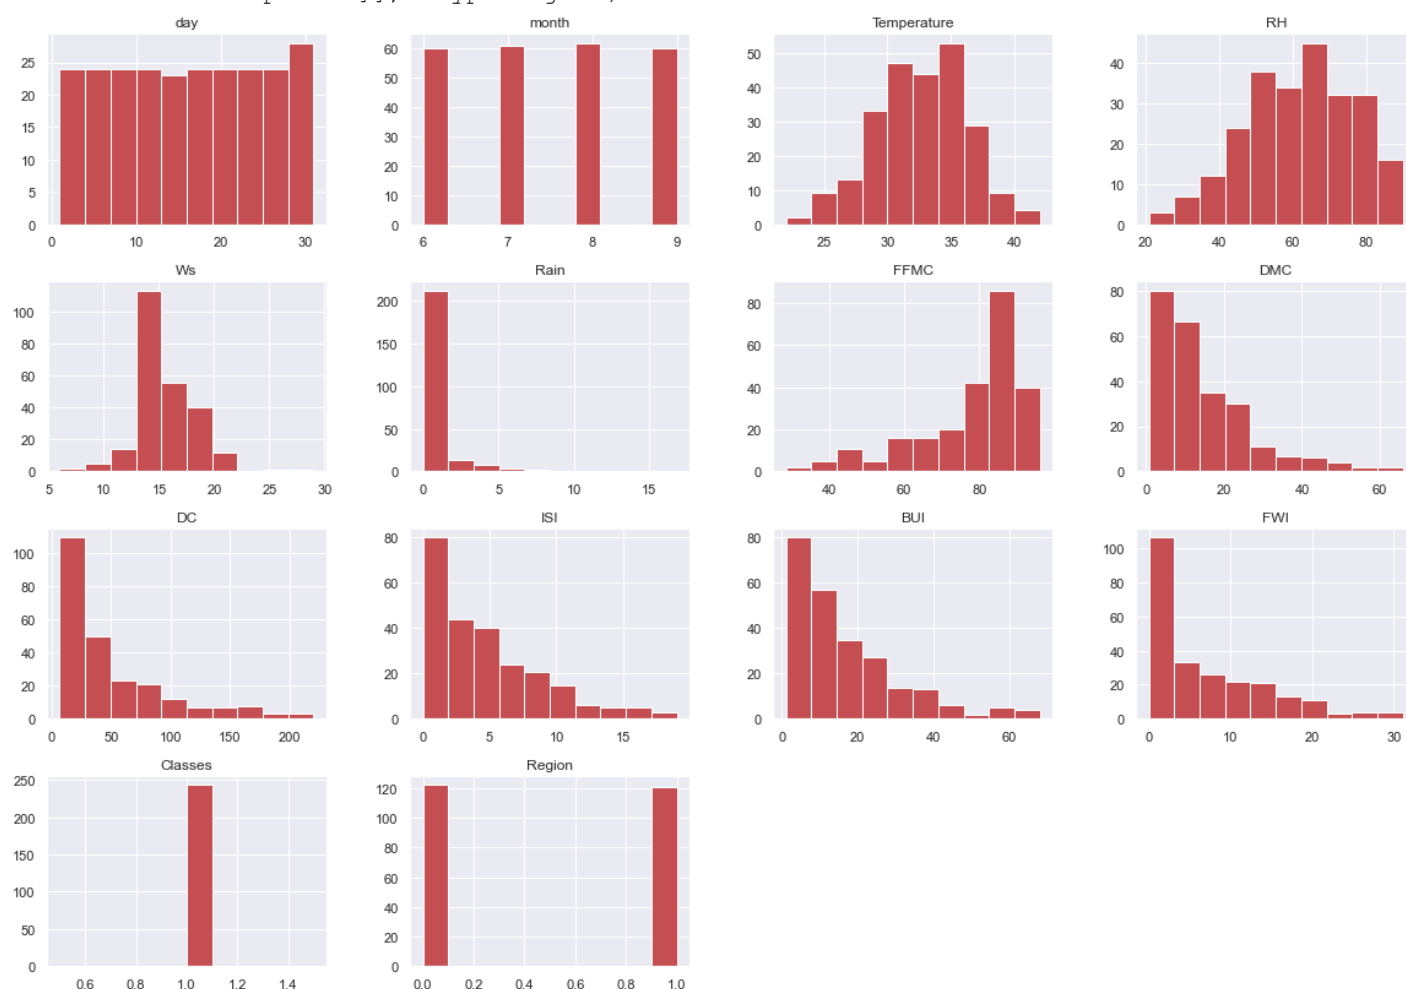


Histogram

-A histogram is basically used to represent data provided in a form of sme groups

```
In [38]: df1.hist(figsize=(20,14),color='r')
```

```
Out[38]: array([[<AxesSubplot:title={'center':'day'}>,  
      <AxesSubplot:title={'center':'month'}>,  
      <AxesSubplot:title={'center':'Temperature'}>,  
      <AxesSubplot:title={'center':'RH'}>],  
      [<AxesSubplot:title={'center':'Ws'}>,  
      <AxesSubplot:title={'center':'Rain'}>,  
      <AxesSubplot:title={'center':'FFMC'}>,  
      <AxesSubplot:title={'center':'DMC'}>],  
      [<AxesSubplot:title={'center':'DC'}>,  
      <AxesSubplot:title={'center':'ISI'}>,  
      <AxesSubplot:title={'center':'BUI'}>,  
      <AxesSubplot:title={'center':'FWI'}>],  
      [<AxesSubplot:title={'center':'Classes'}>,  
      <AxesSubplot:title={'center':'Region'}>], <AxesSubplot:>], dtype=object)
```




Percentage for Pie Chart

```
In [39]: percentage=df.Classes.value_counts(normalize=True)*100  
percentage
```

```
Out[39]: fire          56.378601  
not fire    43.621399  
Name: Classes, dtype: float64
```

Plotting Pie chart


```
In [40]: classes_labels=['Fire','Not Fire']  
plt.figure(figsize=(15,10))  
plt.pie(percentage,labels=classes_labels,autopct="%1.1f%%")  
plt.title("Pie Chart of Classes",fontsize=15)  
plt.show()
```



Forest Fire Temperature Box Plot

```
In [41]: sns.boxplot(df1['Temperature'])
```


```
Out[41]: <AxesSubplot:xlabel='Temperature'>
```



Forest Fire Weather Index System[FWI] Boxplot

```
In [42]: sns.boxplot(df['FWI'],color='orange')
```

```
Out[42]: <AxesSubplot:xlabel='FWI'>
```

4. Building Linear Regression Model

In [43]: `df1`

Out[43]:

	day	month	Temperature	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	Region
0	1	6	29	57	18	0.0	65.7	3.4	7.6	1.3	3.4	0.5	1	0
1	2	6	29	61	13	1.3	64.4	4.1	7.6	1.0	3.9	0.4	1	0
2	3	6	26	82	22	13.1	47.1	2.5	7.1	0.3	2.7	0.1	1	0
3	4	6	25	89	13	2.5	28.6	1.3	6.9	0.0	1.7	0.0	1	0
4	5	6	27	77	16	0.0	64.8	3.0	14.2	1.2	3.9	0.5	1	0
...
239	26	9	30	65	14	0.0	85.4	16.0	44.5	4.5	16.9	6.5	1	1
240	27	9	28	87	15	4.4	41.1	6.5	8.0	0.1	6.2	0.0	1	1
241	28	9	27	87	29	0.5	45.9	3.5	7.9	0.4	3.4	0.2	1	1
242	29	9	24	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7	1	1
243	30	9	24	64	15	0.2	67.3	3.8	16.5	1.2	4.8	0.5	1	1

243 rows × 14 columns

In [44]: `x=df1.drop('Temperature',axis=1)`

```
In [45]: y=df1['Temperature']
```

```
In [46]: x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.33, random_state=10)
```

```
In [47]: x_train
```

Out[47]:

	day	month	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	Region	
	25	26	6	64	18	0.0	86.8	17.8	71.8	6.7	21.6	10.6	1	0
	121	30	9	78	14	1.4	45.0	1.9	7.5	0.2	2.4	0.1	1	0
	174	23	7	71	17	0.0	87.3	46.6	99.0	6.9	46.5	16.3	1	1
	72	12	8	51	13	0.3	81.3	15.6	75.1	2.5	20.7	4.2	1	0
	186	4	8	35	15	0.0	93.8	23.0	42.7	15.7	22.9	20.9	1	1

	64	4	8	69	13	0.0	85.0	8.2	19.8	4.0	8.2	3.9	1	0
	15	16	6	89	13	0.7	36.1	1.7	7.6	0.0	2.2	0.0	1	0
	229	16	9	26	13	0.0	93.9	21.2	59.2	14.2	22.4	19.3	1	1
	125	4	6	64	14	0.0	79.4	5.2	15.4	2.2	5.6	1.0	1	1
	9	10	6	79	12	0.0	73.2	9.5	46.3	1.3	12.6	0.9	1	0

162 rows × 13 columns

```
In [48]: x_test
```

Out[48]:

	day	month	RH	Ws	Rain	FFMC	DMC	DC	ISI	BUI	FWI	Classes	Region	
	46	17	7	70	14	0.0	82.8	9.4	34.1	3.2	11.1	3.6	1	0
	226	13	9	49	19	0.0	88.6	11.5	33.4	9.1	12.4	10.3	1	1
	181	30	7	56	16	0.0	88.9	23.8	57.1	8.2	23.8	13.2	1	1
	116	25	9	81	21	5.8	48.6	3.0	7.7	0.4	3.0	0.1	1	0
	124	3	6	80	14	2.0	48.7	2.2	7.6	0.3	2.6	0.1	1	1

	127	6	6	54	11	0.1	83.7	8.4	26.3	3.1	9.3	3.1	1	1
	242	29	9	54	18	0.1	79.7	4.3	15.2	1.7	5.1	0.7	1	1
	208	26	8	37	16	0.0	92.2	61.3	167.2	13.1	64.0	30.3	1	1
	102	11	9	77	21	1.8	58.5	1.9	8.4	1.1	2.4	0.3	1	0
	78	18	8	54	18	0.0	89.4	20.0	110.9	9.7	27.5	16.1	1	0

81 rows × 13 columns

```
In [49]: y_train
```

```
Out[49]: 25      31
          121     25
          174     31
          72      35
          186     38
          ..
          64      34
          15      29
          229     33
          125     30
          9       28
Name: Temperature, Length: 162, dtype: int32
```

```
In [50]: y_test
```

```
Out[50]: 46      29
          226     29
          181     36
          116     26
          124     29
          ..
          127     35
          242     24
          208     33
          102     30
          78      36
Name: Temperature, Length: 81, dtype: int32
```

```
In [51]: scaler=StandardScaler()
```

```
In [52]: x_train_scaled=scaler.fit_transform(X_train)
```

```
In [53]: x_train_scaled
```

```
Out[53]: array([[ 1.30705791, -1.39305207,  0.06835876, ...,  0.52024214,
                  0.          , -1.01242284],
                [ 1.77217242,  1.29354835,  0.99672801, ..., -0.93452011,
                  0.          , -1.01242284],
                [ 0.95822202, -0.4975186 ,  0.53254338, ...,  1.30997022,
                  0.          ,  0.9877296 ],
                ...,
                [ 0.14427163,  1.29354835, -2.45150064, ...,  1.72561657,
                  0.          ,  0.9877296 ],
                [-1.2510719 , -1.39305207,  0.06835876, ..., -0.8098262 ,
                  0.          ,  0.9877296 ],
                [-0.55340014, -1.39305207,  1.0630401 , ..., -0.82368108,
                  0.          , -1.01242284]])
```

```
In [54]: x_test_scaled=scaler.transform(X_test)
```

```
In [55]: x_test_scaled
```

```
Out[55]: array([[ 0.26055026, -0.4975186 ,  0.46623129, ..., -0.44959936,
                  0.          , -1.01242284],
                [-0.20456425,  1.29354835, -0.92632258, ...,  0.4786775 ,
                  0.          ,  0.9877296 ],
                [ 1.77217242, -0.4975186 , -0.46213796, ...,  0.88046898,
                  0.          ,  0.9877296 ],
                ...,
                ...])
```

```
[ 1.30705791,  0.39801488, -1.72206765, ...,  3.24965322,
  0.          ,  0.9877296 ],
[-0.43712151,  1.29354835,  0.93041592, ..., -0.90681035,
  0.          , -1.01242284],
[ 0.37682889,  0.39801488, -0.59476213, ...,  1.28226046,
  0.          , -1.01242284]])
```

Model Training

```
In [56]: regression=LinearRegression()
```

```
In [57]: regression.fit(x_train_scaled,y_train)
```

```
Out[57]: ▾ LinearRegression
LinearRegression()
```

```
In [58]: #print the coefficients & intercept
print(regression.coef_)
#regression.score(x_test_scaled,y_test)

[-0.36440793 -0.22635361 -1.44134283 -0.7142619  -0.23278709  0.96421264
 -0.01542755  0.65162217  0.13076922  0.26277398 -0.28088408  0.
  0.24272792]
```

```
In [59]: print(regression.intercept_)
```

```
32.074074074074076
```

```
In [60]: reg_pred=regression.predict(x_test_scaled)
```


```
In [61]: reg_pred
```

```
Out[61]: array([31.84063173, 33.08286672, 33.29508782, 24.75854227, 29.24684893,
 33.62257975, 31.75352242, 34.58358479, 31.72328528, 32.36866716,
 33.67767751, 33.25839479, 35.65738396, 31.97503396, 34.04296478,
 33.22479669, 26.47624061, 36.04714436, 33.19344089, 23.64162513,
 32.1343707 , 32.60093493, 33.02495519, 32.94209422, 29.99121682,
 32.68893888, 33.04332061, 32.27268269, 32.09347766, 34.07041738,
 34.51057574, 33.74628677, 34.33569821, 32.69179752, 31.163211 ,
 28.76146317, 32.70115088, 31.76403503, 33.04719502, 34.07436881,
 33.80864285, 35.5411224 , 34.22166608, 37.18514799, 32.83625502,
 36.64474361, 32.38777677, 35.33018647, 30.75190986, 30.94859328,
 32.28439034, 39.1982583 , 32.58656685, 34.55522617, 27.14324627,
 36.9254135 , 33.79514082, 33.90475288, 29.13300329, 32.27931989,
 32.48201446, 32.01342537, 24.82170293, 36.32093954, 36.06160239,
 29.12469145, 29.4759838 , 29.4409954 , 36.02022222, 28.49852589,
 29.20955152, 32.15466961, 30.3678962 , 30.98622517, 34.80797612,
 32.73436681, 35.23626121, 31.37326569, 37.001972 , 26.92983473,
 33.55759223])
```

ASSUMPTIONS OF LINEAR REGRESSION

```
In [62]: #if we are plotting y_test,reg_pred relation should be linear
plt.scatter(y_test,reg_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")
```


Out[62]: Text(0, 0.5, 'Test Predicted Data')



```
In [63]: residuals=y_test-reg_pred
```


```
In [64]: sns.displot(residuals,kind="kde")
```

Out[64]: <seaborn.axisgrid.FacetGrid at 0x1e804eb1130>



```
In [65]: plt.scatter(reg_pred,residuals)
```

Out[65]: <matplotlib.collections.PathCollection at 0x1e806189ca0>



Performance Metrics

```
In [66]: print(mean_squared_error(y_test,reg_pred))
print(mean_absolute_error(y_test,reg_pred))
print(np.sqrt(mean_squared_error(y_test,reg_pred)))
```

```
7.718372490369679
2.1359729920427832
2.7781959056858603
```

R SQUARE & ADJUSTED R SQUARE

```
In [67]: score=r2_score(y_test,reg_pred)
score
```

Out[67]: 0.442435458586767

Now we test the accuracy of the model using Adjusted R2

But why to test the accuracy of the model using Adjusted R2,when we have already tested the accuracy of the model using R2,

It is because the accuracy of the model under R2 will increase as and when we increase the no of features(x)

But the Adjusted R2 will remain unaffected inspite of No of features(x) we take , so R2 is less reliable then Adjusted R2

```
In [68]: ## Adjusted R square
```

```
#display adjusted R-squared
1 - (1-score)*(len(y_test)-1)/(len(y_test)-x_test_scaled.shape[1]-1)
```

Out[68]: 0.3342512938349457

5. Ridge Regression

Ridge Regression (also called as L2 Regularization), the main aim of ridge regression is to reduce overfitting

In [109... `ridge=Ridge()`

In [70]: `ridge.fit(x_train_scaled,y_train)`

Out[70]:

▼ Ridge

Ridge()

In [71]: `print(ridge.coef_)`

```
[-0.35859242 -0.22361881 -1.4246319  -0.71357223 -0.23419537  0.96377519
 0.02101355  0.62937949  0.10421058  0.21661456 -0.22339104  0.
 0.23878478]
```

In [72]: `print(ridge.intercept_)`

```
32.074074074074076
```

In [73]: `print(ridge.intercept_)`

```
32.074074074074076
```

In [74]: `ridge_pred=ridge.predict(x_test_scaled)`


In [75]: `ridge_pred`

Out[75]: `array([31.84522025, 33.07801866, 33.32143372, 24.79311009, 29.25112242,`
33.64529216, 31.73250629, 34.61248132, 31.74324356, 32.36616109,
33.66795511, 33.25653199, 35.63844353, 31.96899817, 33.96914585,
33.23538038, 26.51482354, 36.03947803, 33.1811996 , 23.66716673,
32.09793728, 32.58109242, 33.0232019 , 32.94038806, 29.98560291,
32.70770994, 33.04393524, 32.25760392, 32.08900203, 34.04925765,
34.50009163, 33.73910985, 34.31012523, 32.69207555, 31.13817866,
28.77144618, 32.67740941, 31.75746455, 33.04099284, 34.04747401,
33.83814527, 35.52736647, 34.22117957, 37.13751492, 32.83627381,
36.68161808, 32.36960708, 35.31854183, 30.73882088, 30.95562496,
32.291368 , 39.20825907, 32.57797449, 34.55202448, 27.17475543,
36.9522104 , 33.81487792, 33.86236693, 29.15322938, 32.26413039,
32.47914425, 32.00320521, 24.82974268, 36.34245785, 36.10298328,
29.15073976, 29.45086263, 29.43904153, 36.0418111 , 28.50689289,
29.24000588, 32.16520447, 30.37216231, 30.99205138, 34.80964697,
32.7546397 , 35.207702 , 31.36583658, 37.04199234, 26.94898166,
33.54871979])

ASSUMPTIONS OF LINEAR REGRESSION

```
In [76]: #if we are plotting y_test, reg_pred relation should be linear  
plt.scatter(y_test, ridge_pred)  
plt.xlabel("Test Truth Data")  
plt.ylabel("Test Predicted Data")
```


Out[76]: Text(0, 0.5, 'Test Predicted Data')



```
In [77]: residuals=y_test-ridge_pred
```


```
In [78]: sns.displot(residuals, kind="kde")
```

Out[78]: <seaborn.axisgrid.FacetGrid at 0x1e8061e60a0>



```
In [79]: plt.scatter(ridge_pred, residuals)
```

```
Out[79]: <matplotlib.collections.PathCollection at 0x1e8066e2400>
```



Performance Metrics

```
In [80]: print(mean_squared_error(y_test, ridge_pred))  
print(mean_absolute_error(y_test, ridge_pred))  
print(np.sqrt(mean_squared_error(y_test, ridge_pred)))
```

7.700169527845219

2.1363570878417484
2.774917931731535

R SQUARE & ADJUSTED R SQUARE

```
In [81]: score=r2_score(y_test,reg_pred)
score
```

Out[81]: 0.442435458586767

Now we test the accuracy of the model using Adjusted R2

But why to test the accuracy of the model using Adjusted R2,when we have already tested the accuracy of the model using R2,

It is because the accuracy of the model under R2 will increase as and when we increase the no of features(x)

But the Adjusted R2 will remain unaffected inspite of No of features(x) we take , so R2 is less reliable then Adjusted R2

```
In [82]: ## Adjusted R square
#display adjusted R-squared
1 - (1-score)*(len(y_test)-1)/(len(y_test)-x_test_scaled.shape[1]-1)
```

Out[82]: 0.3342512938349457

6. LASSO

- LASSO Regression(also called as L1 Norm/L1 Regularization),the main aim of LASSO regressio is to reduce the features which are least correlated with dependent variables .
- If our datasets has outliers we should use LASSO .

```
In [83]: lasso=Lasso()
lasso
```

Out[83]: ▼ Lasso
Lasso()

```
In [84]: lasso.fit(x_train_scaled,y_train)
```

Out[84]: ▼ Lasso
Lasso()

```
In [85]: print(lasso.coef_)
```

```
[ 0.          -0.          -0.88423537 -0.          -0.          0.88313134
 0.           0.           0.           0.           0.           0.
 0.]
```

```
In [86]: print(lasso.intercept_)
```

```
32.074074074074076
```

```
In [87]: lasso_pred=lasso.predict(x_test_scaled)
```


```
In [88]: lasso_pred
```

```
Out[88]: array([31.99263189, 33.56261409, 33.16968133, 29.35085197, 29.41532603,
                33.00570181, 32.24618859, 33.57987995, 32.17946599, 33.03180429,
                32.43352237, 32.64830012, 34.42463082, 31.61671515, 32.83763259,
                33.35142638, 29.33942467, 34.79854886, 32.38506498, 29.63719153,
                32.36854863, 32.21265622, 32.60134175, 31.81897394, 31.16330572,
                32.75364414, 32.88409128, 32.36320974, 32.04108928, 31.96868565,
                33.24608236, 33.07067555, 33.54659742, 32.57214892, 31.41161583,
                30.06340682, 32.79726231, 32.13550568, 33.36310351, 33.22497663,
                33.25167109, 34.07840658, 33.22831682, 33.75728546, 32.90619636,
                34.59903828, 32.34644356, 33.72050521, 31.73648451, 31.41845375,
                32.26170559, 35.46055533, 32.15686124, 33.60932262, 29.07618954,
                34.41195434, 33.23999395, 32.50358516, 29.88575148, 32.20406944,
                32.0385909 , 31.44605525, 29.20463799, 34.49394409, 33.4095622 ,
                30.24824223, 31.37474327, 30.84852284, 34.06622977, 30.49062147,
                31.23012062, 32.0385909 , 31.43946717, 31.10541974, 33.38061921,
                32.30257555, 32.9833469 , 32.74980427, 34.4764284 , 30.16341195,
                33.31614515])
```

ASSUMPTIONS OF LINEAR REGRESSION

```
In [89]: #if we are plotting y_test,reg_pred relation should be linear
plt.scatter(y_test,lasso_pred)
plt.xlabel("Test Truth Data")
plt.ylabel("Test Predicted Data")
```


```
Out[89]: Text(0, 0.5, 'Test Predicted Data')
```



```
In [90]: residuals=y_test-lasso_pred
```


```
In [91]: sns.displot(residuals,kind="kde")
```

```
Out[91]: <seaborn.axisgrid.FacetGrid at 0x1e875aaa4c0>
```



```
In [92]: plt.scatter(lasso_pred,residuals)
```

```
Out[92]: <matplotlib.collections.PathCollection at 0x1e806c50ac0>
```



Performance Metrics

```
In [93]: print(mean_squared_error(y_test,lasso_pred))
print(mean_absolute_error(y_test,lasso_pred))
print(np.sqrt(mean_squared_error(y_test,lasso_pred)))
```

```
8.69462464944341
2.370968686727018
2.9486648927003234
```

R SQUARE & ADJUSTED R SQUARE

```
In [94]: score=r2_score(y_test,lasso_pred)
score
```

```
Out[94]: 0.3719123543887275
```

```
In [95]: ## Adjusted R square
#display adjusted R-squared
1 - (1-score)*(len(y_test)-1)/(len(y_test)-x_test_scaled.shape[1]-1)
```

```
Out[95]: 0.250044602255197
```

7. ElasticNet

ElasticNet is a combination of Ridge & LAASO Regression

```
In [96]: elastic_net=ElasticNet()
```

```
In [97]: elastic_net.fit(x_train_scaled,y_train)
```

```
Out[97]: ▼ ElasticNet  
ElasticNet()
```

```
In [98]: print(elastic_net.coef_)
```

```
[-0.          -0.          -0.77155493 -0.27327033 -0.02945645  0.70980198  
 0.11177449  0.          0.20914599  0.04593431  0.12829249  0.  
 0.          ]
```

```
In [99]: print(elastic_net.intercept_)
```

```
32.074074074074076
```

```
In [100... elastic_net_pred=elastic_net.predict(x_test_scaled)
```


```
In [101... elastic_net_pred
```

```
Out[101... array([31.9573208 , 33.23686908, 33.35699564, 28.69175409, 29.50702659,  
      33.21952486, 31.75246849, 33.96237735, 32.10800305, 32.71913539,  
      32.6464949 , 32.70818669, 34.28160866, 31.98757886, 32.58127398,  
      33.38088877, 29.25544946, 35.1329971 , 32.41345187, 28.81942678,  
      31.92062398, 32.01321369, 32.87611371, 32.19248582, 30.76147509,  
      32.77885483, 32.9241576 , 31.80574265, 31.84088675, 32.63570681,  
      33.35824365, 32.97725238, 33.16526411, 32.6259555 , 31.48228091,  
      29.90876877, 32.28704956, 31.73349961, 33.28556103, 32.87586141,  
      33.61368581, 34.06687234, 33.48616338, 34.31378768, 32.80354959,  
      35.56194932, 31.98083723, 34.20080511, 31.06352775, 31.3703163 ,  
      32.33015611, 36.68465163, 32.02910227, 33.90164802, 29.13016242,  
      35.33664404, 33.44975962, 32.59500157, 29.89482886, 32.07728159,  
      32.39209933, 31.39172713, 28.33121558, 35.11153428, 34.62534381,  
      30.33956758, 30.58992082, 30.41863187, 34.79501455, 29.86618003,  
      31.0520072 , 32.09680635, 31.20338706, 31.12722538, 34.01998343,  
      32.49121998, 33.0901132 , 32.02987963, 35.52196854, 29.45679071,  
      33.38350807])
```

ASSUMPTIONS OF LINEAR REGRESSION

```
In [102... #if we are plotting y_test,reg_pred relation should be linear  
plt.scatter(y_test,elastic_net_pred)  
plt.xlabel("Test Truth Data")  
plt.ylabel("Test Predicted Data")
```


```
Out[102... Text(0, 0.5, 'Test Predicted Data')
```



```
In [103... residuals=y_test-elastic_net_pred
```


```
In [104... sns.displot(residuals,kind="kde")
```

```
Out[104... <seaborn.axisgrid.FacetGrid at 0x1e806949280>
```



```
In [105... plt.scatter(elastic_net_pred,residuals)
```

```
Out[105... <matplotlib.collections.PathCollection at 0x1e80715d2b0>
```



Performance Metrics

In [106...

```
print(mean_squared_error(y_test,elastic_net_pred))
print(mean_absolute_error(y_test,elastic_net_pred))
print(np.sqrt(mean_squared_error(y_test,elastic_net_pred)))
```

```
7.918267418307838
2.2796482511865195
2.813941616009088
```

R SQUARE & ADJUSTED R SQUARE

In [107...

```
score=r2_score(y_test,elastic_net_pred)
score
```

Out[107...

```
0.4279953257782334
```

Now we test the accuracy of the model using Adjusted R2

But why to test the accuracy of the model using Adjusted R2,when we have already tested the accuracy of the model using R2,

It is because the accuracy of the model under R2 will increase as and when we increase the no of features(x)

But the Adjusted R2 will remain unaffected inspite of No of features(x) we take , so R2 is less reliable then Adjusted R2

In [108...

```
## Adjusted R square
#display adjusted R-squared
1 - (1-score)*(len(y_test)-1)/(len(y_test)-x_test_scaled.shape[1]-1)
```


0.317009344212816

Out[108...